

```
In [1]: import itertools
import numpy as np
import pandas as pd
from pprint import pprint
import matplotlib.pyplot as plt

from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import KFold, cross_validate, train_test_split
from sklearn.metrics import silhouette_score, mean_absolute_error, mean_squared_error

from sklearn.svm import SVR
from sklearn.tree import DecisionTreeRegressor
from sklearn.neural_network import MLPRegressor
from sklearn.neighbors import KNeighborsRegressor
from sklearn.linear_model import LogisticRegression, Lasso, BayesianRidge, LinearRegression

from sklearn.cluster import KMeans
```

```
In [2]: ds_1 = pd.read_csv('ds1.csv', index_col = 0)
ds_2 = pd.read_csv('ds2.csv', index_col = 0)
print(ds_1.isnull().values.any())
print(ds_2.isnull().values.any())
```

```
False
False
```

```
In [3]: ds_1.head()
```

```
Out[3]:
```

| | x1 | x2 | x3 | x5 | x6 | ya | yb | yc |
|---|----------|-----------|----------|----------|-----------|-----------|----------|-----------|
| 1 | 2.642583 | -1.715220 | 1.909334 | 0.027139 | -3.447187 | 13.630850 | 1.828765 | 0.008386 |
| 2 | 4.588761 | -2.507543 | 4.239107 | 1.704150 | -2.782809 | 7.834582 | 2.162110 | 0.000008 |
| 3 | 7.919796 | -5.108415 | 3.039451 | 0.992815 | 5.551587 | -5.107041 | 2.797083 | -0.000005 |
| 4 | 2.616757 | -2.124040 | 2.855570 | 0.990079 | 1.694697 | 19.015046 | 1.953887 | 0.038017 |
| 5 | 3.300856 | -5.159684 | 0.764544 | 0.143581 | 3.277496 | -9.818862 | 1.922446 | 0.001178 |

```
In [4]: ds_2.head()
```

```
Out[4]:
```

| | X1 | X2 | X3 | X4 | X5 | X6 | X7 | X8 |
|---|-----------|-----------|-----------|------------|-----------|------------|-----------|------------|
| 1 | 23.778224 | 13.319974 | 15.565124 | -3.713626 | 7.296793 | -19.371013 | -0.894130 | -6.110282 |
| 2 | 16.602950 | 23.311281 | 21.099052 | -0.304154 | -3.218990 | 2.357643 | 12.027277 | 7.070349 |
| 3 | 12.084683 | 19.710443 | 9.837102 | -1.081918 | -1.201942 | 9.738019 | 16.125920 | 19.119391 |
| 4 | 13.044534 | 10.749040 | 5.884407 | -11.703525 | -4.134358 | -22.344666 | -1.263349 | 0.493711 |
| 5 | 8.314115 | 6.748794 | 5.388535 | -0.000290 | -4.724787 | -16.346812 | 3.293600 | -10.848273 |

Question 1

a. Describe the data set in a few sentences.

E.g. What are the distributions of each feature?

Summary statistics?

b. Try to come up with a predictive model, e.g. $y = f(x_1, \dots, x_n)$ for each y sequence.

Describe your models and how you came up with them.

What (if any) are the predictive variables?

How good would you say each of your models is?

a)

The data set has 5 features (the x columns) and 3 targets (the y columns).

On a side note, not sure if the missing `x4` is a typo or intentional, but thats why I decided to create the 2 functions `get_targets_list` and `get_features_list` to get the `x` and `y` columns from the dataset.

Class `Describe` does the EDA for the data. Each feature column contains float values. None of the columns contain any null values, if they did, I would use the mean for each column to fill them, using `df.fillna(df.mean())`

I plot the histogram and the kernel density estimation (kde) plots for each feature.

For summary stats, I list out the count, mean, standard deviation, min, max, 25% quantile, 50% quantile, 75% quantile, median, skew, average (same as mean), variance, sum of all values and standard error of the mean (sem) for each of the features.

Every column contains 10000 values. The values for the summary stas are printed when the `runner` function in `Describe` class is called.

Based on the histogram plots and kde, heres what I can conclude about the features in terms of distribution:

1. **x1** has a Uniform Distribution
2. **x2** has a Cauchy Distribution (could also be Normal Distribution)
3. **x3** has a Normal Distribution
4. **x5** has a F Distribution
5. **x6** has a Double Exponential Distribution

I also made 15 additional plots, one for each individual feature vs each individual target. This can show which features are important for each target, also how the data is scattered, where the data is centered and if there are any outliers.

b) Looking at the data, we can see that its a Regression problem.

I used the basic ML training methodology.

First, I broke the data into a train test split, where 20% of data is for testing.

Next, I used K Fold cross validation (5 splits) to ensure the accuracy of the models, and to

make sure there isn't over fitting.

Then to see what features are useful, and what features are not, I used a permutation combination to get all possible combinations of the features.

I ran each combination through k fold, and recorded the accuracy of the models. See self.scoring array to see all the different metrics I use to measure the performance of the models.

Since this is a regression problem, I tried several different models from SkLearn. The list of models are shown in the first cell up top. The results were very similar, so ultimately I choose `LinearRegression` model.

The `evaluate` function in `Regression_Evaluation` class evaluates the different combinations. And then produces a table of R2 scores for each combination. I only show the top 3 combination sorted by score.

For each target, I then take the top combination of variables, and train a model, and then test it using a train test split of 0.9.

I also use all the features and do the same for each target.

The results are shown below, as well as the exact equation, which is a combination of coefficients and intercept.

I also plot the predicted y values vs the actual y values to show how well the model predicts values, and how close the predicted values are to the true values.

For ya, I get an R2 score of 67%, which shows the model is ok, but needs improvement. This shows that there is a correlation between the features and the target, but it is not strong. For ya, the important features are [x1, x2, x3].

For yb, I get an R3 score of 89%, which shows the model is very good. This shows that there is a strong correlation between the features and the target. For yb, the important features are [x1, x3].

For yc, I get an R3 score of 0%, which shows the model is not at all correct. This shows that there is no correlation between the features and the target, regardless of the combination of features.

```
In [5]: def get_targets_list(ds_1):
        target_list = []
        for col in ds_1.columns:
            if ('y' in col):
                target_list.append(col)
        return target_list

        def get_features_list(ds_1):
            feature_list = []
            for col in ds_1.columns:
                if ('x' in col):
                    feature_list.append(col)
            return feature_list
```

```
In [6]: class Describe:
    def __init__(self, ds_1):
        self.ds_1 = ds_1
        self.kinds = ['hist', 'kde']
        self.summary_list = ['median', 'skew', 'average', 'var', 'sem', 'sum']
        self.features = get_features_list(self.ds_1)

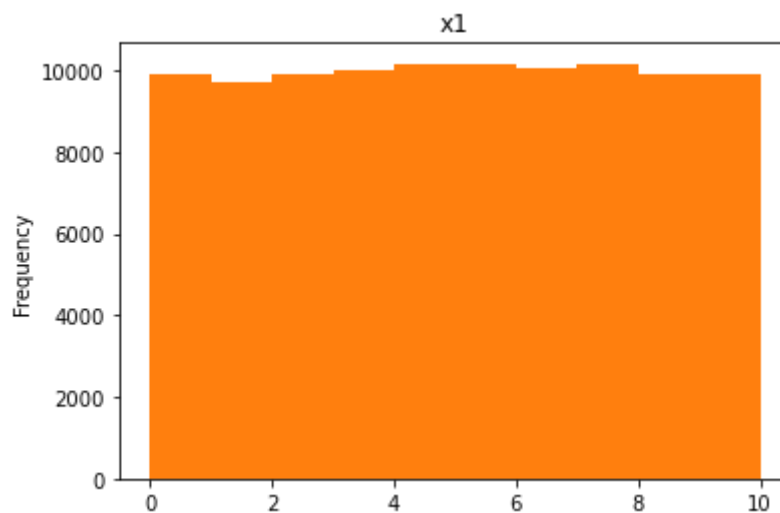
    def distribution(self, feature, title = '', kind = 'hist'):
        feature.plot(kind = kind, title = title)
        feature.plot(kind = kind, title = title)
        plt.show()

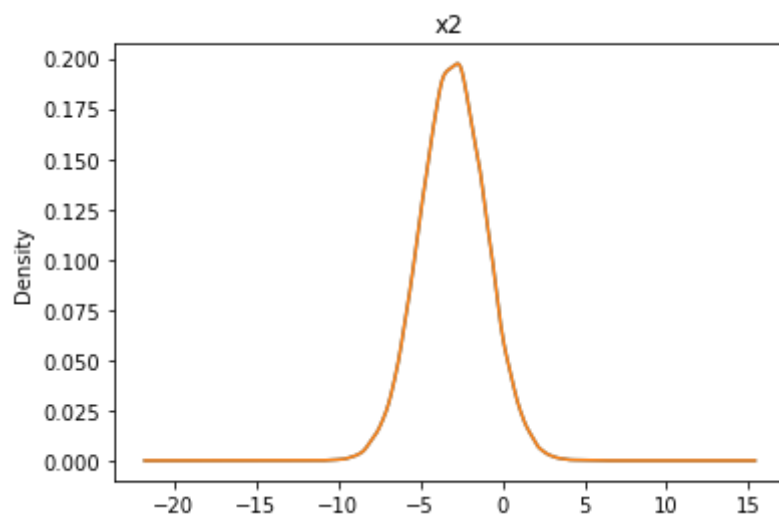
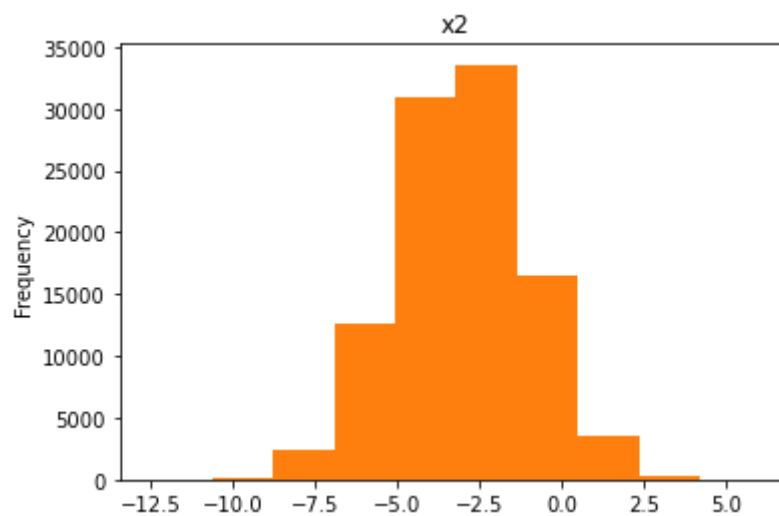
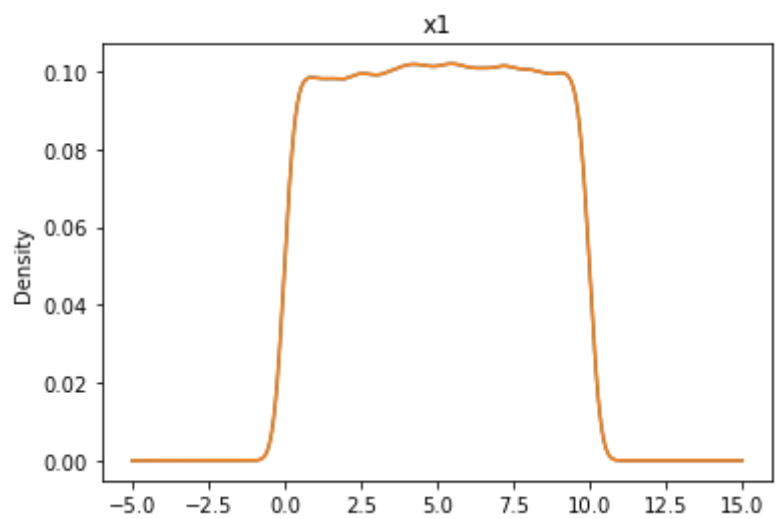
    def summary(self, feature, title = ''):
        print(title, '\n')
        print(feature.describe(), '\n')
        print(feature.agg(self.summary_list), '\n')

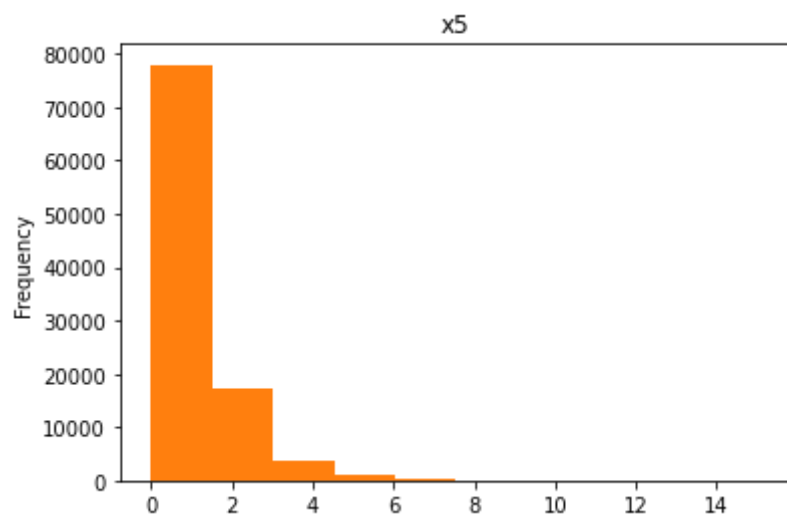
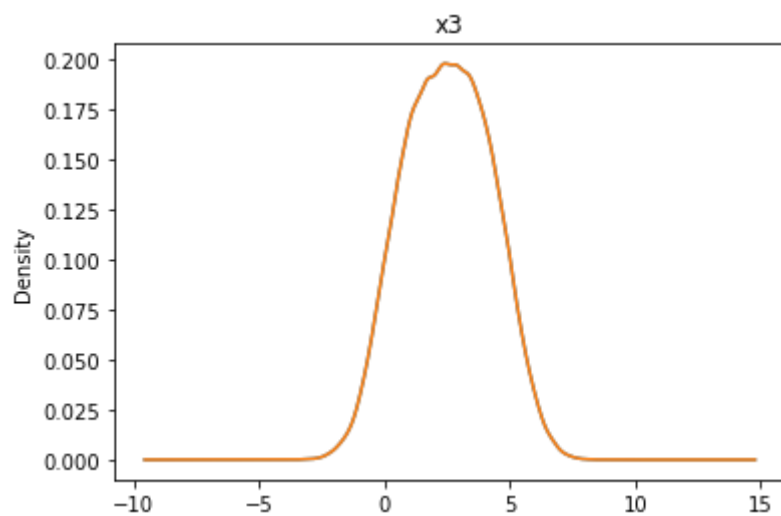
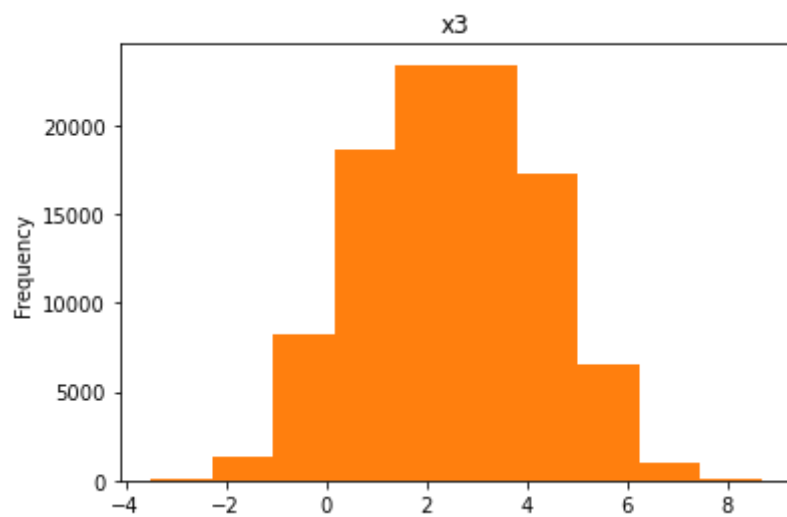
    def runner(self):
        ds_1 = self.ds_1
        for feature in self.features:
            for kind in self.kinds:
                self.distribution(ds_1[feature], title = feature, kind = kind)

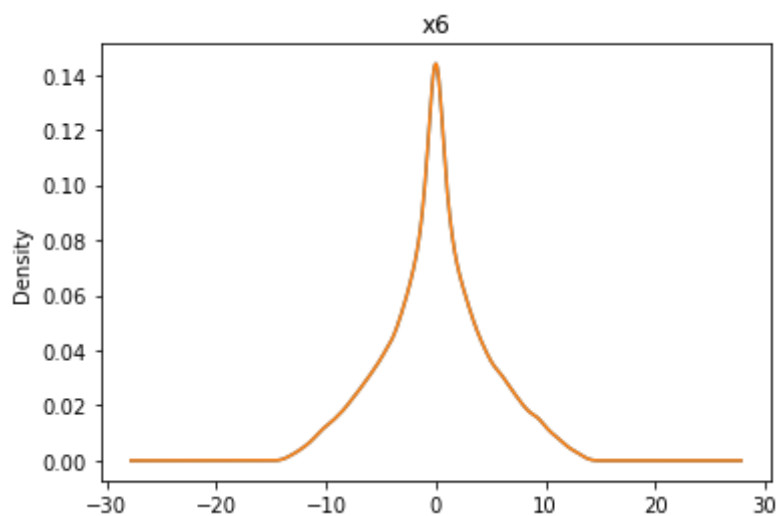
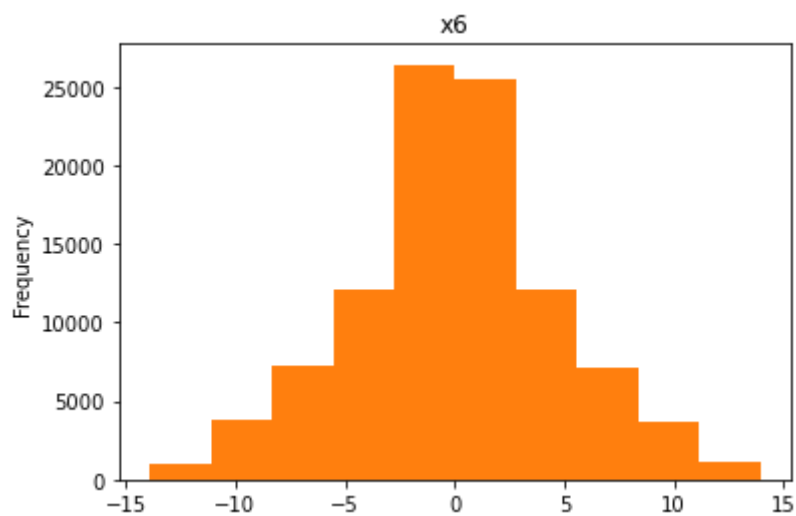
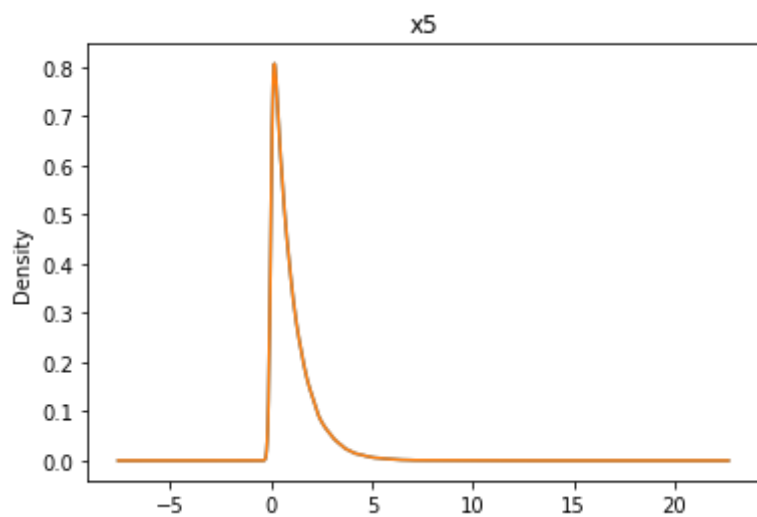
        for feature in self.features:
            self.summary(ds_1[feature], feature)
```

```
In [7]: Describe(ds_1).runner()
```









x1

```
count    100000.000000
mean      5.011059
std       2.873568
min       0.000015
25%       2.536309
50%       5.022191
75%       7.486275
max       9.999887
```

Name: x1, dtype: float64

```
median    5.022191
skew      -0.008450
average    5.011059
var       8.257391
sem       0.009087
sum       501105.931893
```

Name: x1, dtype: float64

x2

```
count    100000.000000
mean     -3.005565
std       2.000799
min      -12.498524
25%      -4.353844
50%      -3.002649
75%      -1.649283
max       6.089820
```

Name: x2, dtype: float64

```
median    -3.002649
skew      -0.000347
average   -3.005565
var       4.003195
sem       0.006327
sum      -300556.464753
```

Name: x2, dtype: float64

x3

```
count    100000.000000
mean      2.500593
std       1.752906
min      -3.489212
25%       1.189920
50%       2.503764
75%       3.802236
max       8.679097
```

Name: x3, dtype: float64

```
median    2.503764
skew      -0.003621
average    2.500593
var       3.072681
sem       0.005543
sum       250059.262904
```

Name: x3, dtype: float64

x5

```

count      100000.000000
mean        0.999136
std         1.002006
min         0.000003
25%         0.285629
50%         0.690903
75%         1.386862
max         15.102966
Name: x5, dtype: float64

```

```

median      0.690903
skew        2.039982
average     0.999136
var         1.004017
sem         0.003169
sum         99913.647275
Name: x5, dtype: float64

```

x6

```

count      100000.000000
mean        0.000647
std         4.663860
min        -13.885453
25%        -2.611943
50%        -0.000611
75%         2.621841
max         13.924740
Name: x6, dtype: float64

```

```

median     -0.000611
skew        0.003579
average     0.000647
var         21.751591
sem         0.014748
sum         64.684123
Name: x6, dtype: float64

```

```

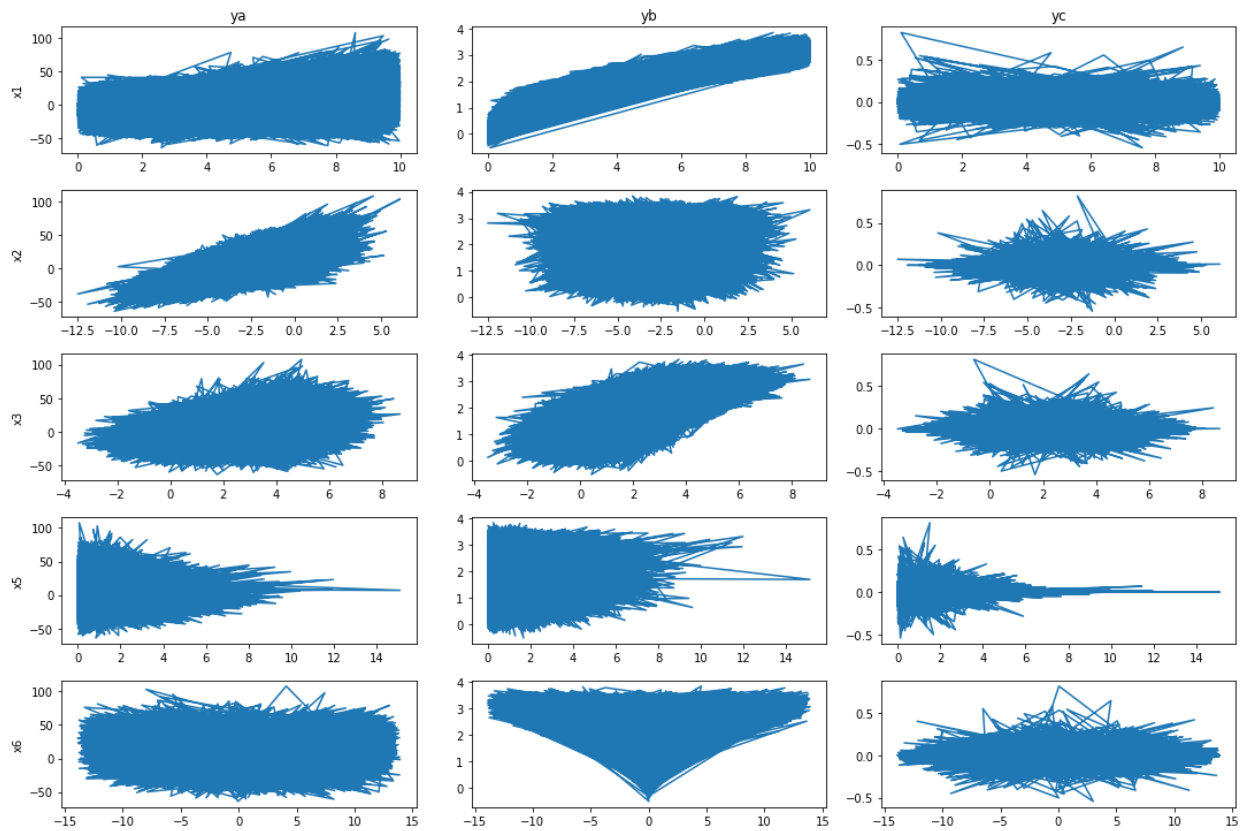
In [8]: def visualize(ds_1):
        targets_list = get_targets_list(ds_1)
        features_list = get_features_list(ds_1)
        fig, ax = plt.subplots(5, 3, figsize = (15, 10))
        fig.tight_layout()
        for i, feature in enumerate(features_list):
            for j, target in enumerate(targets_list):
                ax[i, j].plot(ds_1[feature], ds_1[target])
                ax[0, j].set_title(target)
                ax[i, 0].set_ylabel = feature
        plt.show()

```

```

In [9]: visualize(ds_1)

```



```
In [10]: class Regression_Evaluation:
    def __init__(self, ds_1):
        self.ds_1 = ds_1

        self.model = LinearRegression()
        self.scoring = ['max_error', 'neg_mean_absolute_error',
                        'neg_mean_squared_error', 'neg_root_mean_squared_error', 'r2']

        self.targets = get_targets_list(ds_1)
        self.features = get_features_list(ds_1)
        self.feature_permutations = self.get_feature_permutations(self.features)

    def get_feature_permutations(self, features, subsets = []):
        for length in range(len(features) + 1):
            for subset in itertools.combinations(features, length):
                subsets.append(list(subset))
        return subsets[1:]

    def get_data(self, ds_1, features, target, test_size = 0.2):
        X = ds_1[features]
        y = ds_1[target]
        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = test_size)
        return (X_train, X_test, y_train, y_test)

    def print_score(self, score):
        for key, val in score.items():
            print(key, np.mean(val))

    def kfold(self, X_train, y_train, n_split = 5):
        kf = KFold(n_splits = n_split)
        scores = cross_validate(self.model, X_train, y_train,
                                scoring = self.scoring, cv = kf, return_train_score = True)
        return scores
```

```

def evaluate(self, to_print = False, eval_scores = []):
    feature_permutations = self.feature_permutations
    for target in self.targets:
        for feature_list in feature_permutations:
            X_train, x_test, y_train, y_test = self.get_data(self.ds_1, feature_list)
            scores = self.kfold(X_train, y_train)
            if (to_print):
                print('Target:', target)
                print('Features:', feature_list)
                self.print_score(scores, '\n')
            eval_scores.append({
                'target': target,
                'feature_list': feature_list,
                'score': round(np.mean(scores['test_r2']), 7)
            })
    return pd.DataFrame(eval_scores)

```

In [11]: `scores_features_df = Regression_Evaluation(ds_1).evaluate(to_print = False)`
`scores_features_df.sort_values(['score'], ascending = False).groupby('target').`

Out[11]:

| | target | feature_list | score |
|----|--------|------------------|-----------|
| 37 | yb | [x1, x3] | 0.894749 |
| 31 | yb | [x1] | 0.894744 |
| 49 | yb | [x1, x3, x5] | 0.894744 |
| 15 | ya | [x1, x2, x3] | 0.676735 |
| 25 | ya | [x1, x2, x3, x5] | 0.676721 |
| 26 | ya | [x1, x2, x3, x6] | 0.676717 |
| 65 | yc | [x5] | -0.000037 |
| 72 | yc | [x2, x5] | -0.000046 |
| 63 | yc | [x2] | -0.000047 |

In [12]:

```

def print_equation(features, target, coef, intercept):
    temp = ''
    for i in range(len(features)):
        temp += str(round(coef[i], 4)) + '*' + features[i] + ' '
    return '{} = {} + {}'.format(target, temp[:-1], intercept)

def train(features, target, test_size = 0.1):
    X = ds_1[features]
    y = ds_1[target]
    X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = test_size)
    model = LinearRegression()
    model.fit(X_train, y_train)
    print('Target: {}, Features: {}'.format(target, features))
    print('Model Score (R²):', np.round(model.score(X_test, y_test), 7))
    print('Model Coefficient:', model.coef_)
    print('Model Intercept:', model.intercept_)
    print('Equation:', print_equation(features, target, model.coef_, model.intercept_))

    y_pred = model.predict(X_test)
    plt.scatter(y_pred, y_test)

```

```
plt.title('Predicted vs True Value')
plt.xlabel('Predicted')
plt.ylabel('True')
plt.show()
print()
```

```
best_predictive_variables = scores_features_df.sort_values(['score'], ascending=False)
for index, row in best_predictive_variables.iterrows():
    train(row['feature_list'], row['target'])
    train(get_features_list(ds_1), row['target'])
```

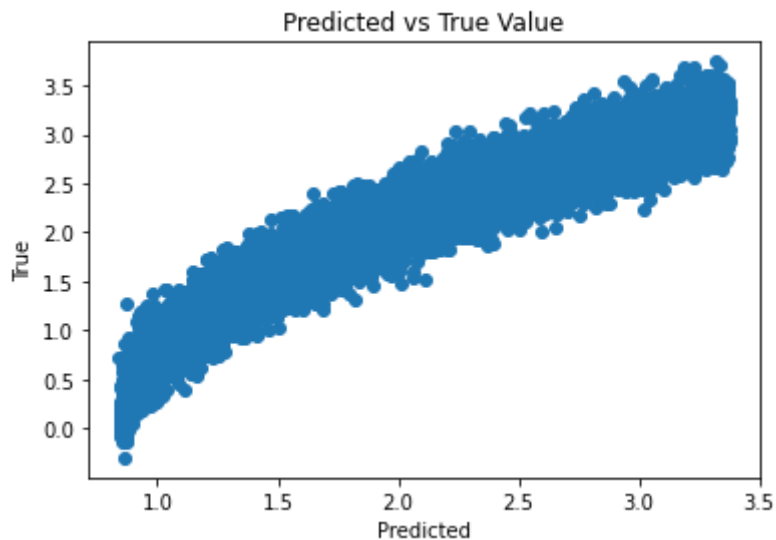
Target: yb, Features: ['x1', 'x3']

Model Score (R^2): 0.8952979

Model Coefficient: [0.25431713 -0.00209127]

Model Intercept: 0.842373986968475

Equation: $y_b = 0.2543 \cdot x_1 - 0.0021 \cdot x_3 + 0.842373986968475$



Target: yb, Features: ['x1', 'x2', 'x3', 'x5', 'x6']

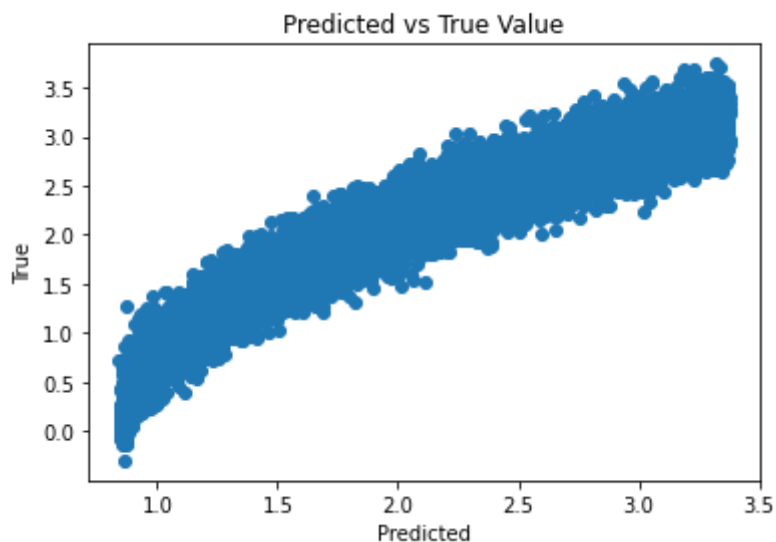
Model Score (R^2): 0.8952963

Model Coefficient: [2.54315379e-01 5.49181362e-05 -2.08919220e-03 4.14899709e-04

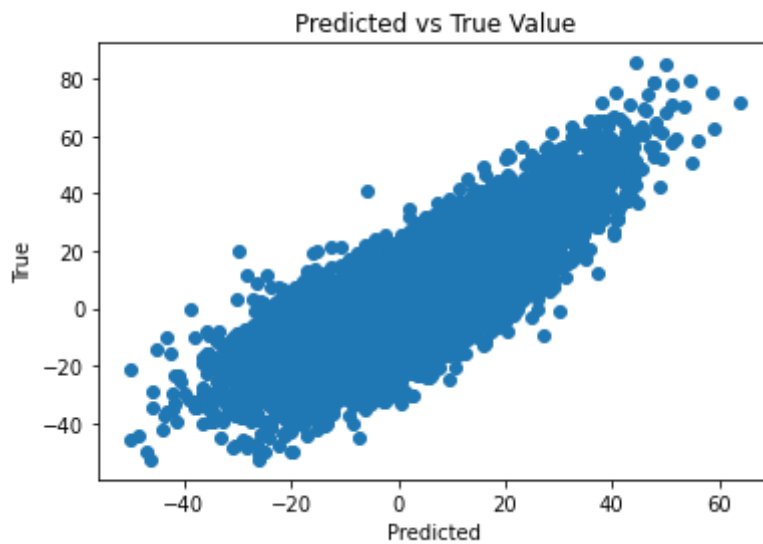
8.24706281e-05]

Model Intercept: 0.8421275974510247

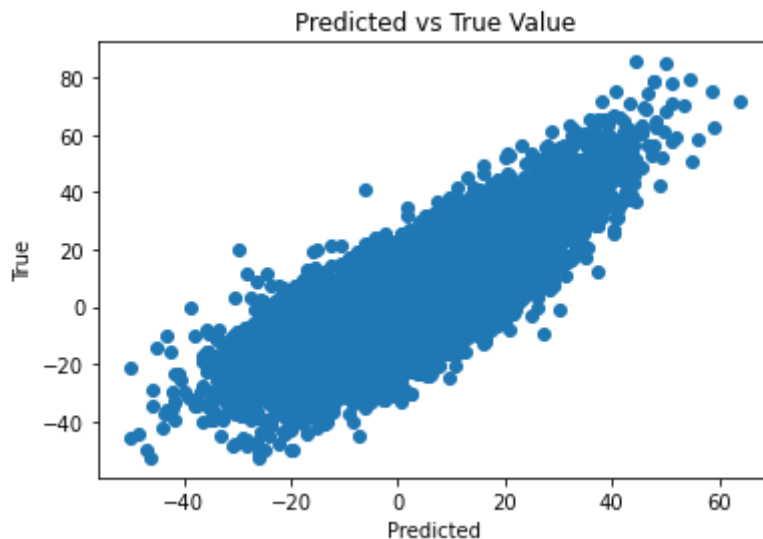
Equation: $y_b = 0.2543 \cdot x_1 + 0.0001 \cdot x_2 - 0.0021 \cdot x_3 + 0.0004 \cdot x_5 + 0.0001 \cdot x_6 + 0.8421275974510247$



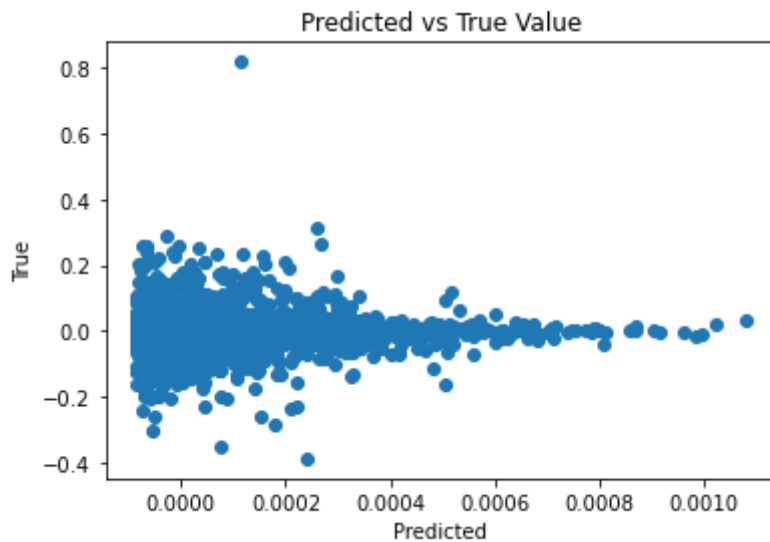
Target: ya, Features: ['x1', 'x2', 'x3']
Model Score (R^2): 0.676793
Model Coefficient: [2.52373744 6.23396718 0.93839982]
Model Intercept: 7.555240400565005
Equation: $ya = 2.5237*x1 + 6.234*x2 + 0.9384*x3 + 7.555240400565005$



Target: ya, Features: ['x1', 'x2', 'x3', 'x5', 'x6']
Model Score (R^2): 0.6767721
Model Coefficient: [2.52368961 6.23392286 0.93844665 0.00952342 0.00653299]
Model Intercept: 7.545714130783381
Equation: $ya = 2.5237*x1 + 6.2339*x2 + 0.9384*x3 + 0.0095*x5 + 0.0065*x6 + 7.545714130783381$



Target: yc, Features: ['x5']
Model Score (R^2): -0.0004263
Model Coefficient: [0.00013382]
Model Intercept: -8.692217632199213e-05
Equation: $yc = 0.0001*x5 - 8.692217632199213e-05$



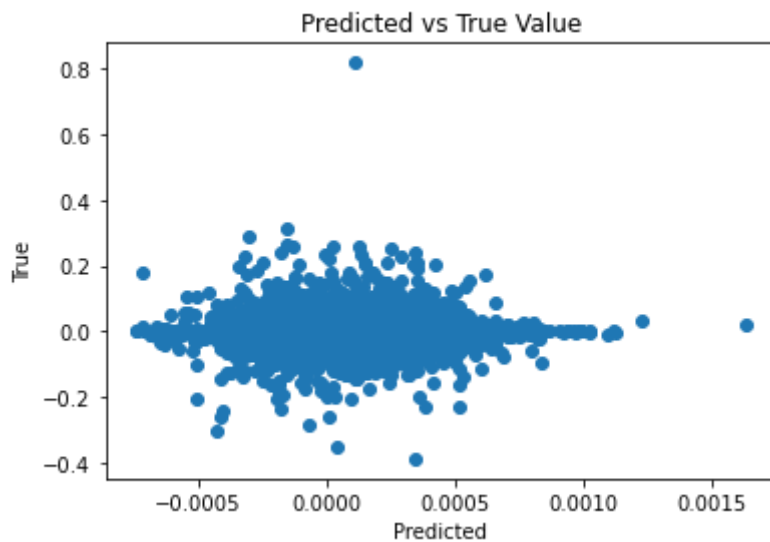
Target: yc, Features: ['x1', 'x2', 'x3', 'x5', 'x6']

Model Score (R^2): -0.0004205

Model Coefficient: [1.24223052e-04 1.09113732e-08 -1.95718882e-04 1.32573583e-04 -8.50130073e-07]

Model Intercept: -0.00021882072471436637

Equation: $yc = 0.0001 \cdot x1 + 0.0 \cdot x2 - 0.0002 \cdot x3 + 0.0001 \cdot x5 - 0.0 \cdot x6 + -0.00021882072471436637$



Question 2

- Describe the data in a few sentences
- How would you visualize this data set?
- Can you identify the number of groups in the data and assign each row to its group?
- Can you create a good visualization of your groupings?

a) I use similar techniques as the first part to describe the data. The results for each column are shown below.

b) This is an unsupervised dataset. There are 10 features. In order to visualize the data, we need to perform PCA on it to reduce the dimensionality to 2. Doing so gives us the graph that is shown below

c) Looking at the output from part b, we can see that there are 4 groups or clusters. But to verify it, I get the silhouette score of the data set. I normalize the data using standard scaler. Then I set a list of clusters, and plot the score of the data for each cluster. This indeed confirms that there are 4 clusters that the data can be grouped into.

Once I do that, I run KMeans Clustering on the data with the number of clusters set to 4, and then I print out the labels, which are the groups for each row.

d) The different coloured clusters and their centers are shown in the plot below

```
In [13]: ds_2.agg(['median', 'skew', 'average', 'var', 'sem', 'sum'])
```

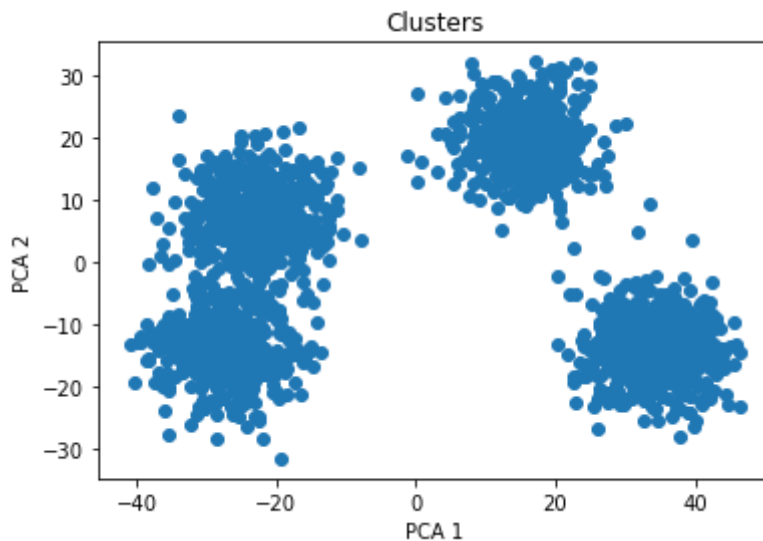
| | X1 | X2 | X3 | X4 | X5 | X6 |
|----------------|--------------|--------------|--------------|--------------|-------------|------------|
| median | 12.754335 | 11.896021 | 11.422441 | -2.631318 | 2.484416 | 1.500838 |
| skew | -0.810685 | -0.106714 | -0.590964 | -0.004560 | 0.121531 | -0.097249 |
| average | 8.677829 | 11.716801 | 9.252817 | -2.679634 | 2.774942 | 0.077631 |
| var | 143.324648 | 44.293462 | 97.248049 | 112.894116 | 77.449728 | 236.405310 |
| sem | 0.267698 | 0.148818 | 0.220509 | 0.237586 | 0.196786 | 0.343806 |
| sum | 17355.657407 | 23433.602030 | 18505.634865 | -5359.268003 | 5549.883133 | 155.261180 |

```
In [14]: ds_2.describe()
```

| | X1 | X2 | X3 | X4 | X5 | X6 |
|--------------|-------------|-------------|-------------|-------------|-------------|-------------|
| count | 2000.000000 | 2000.000000 | 2000.000000 | 2000.000000 | 2000.000000 | 2000.000000 |
| mean | 8.677829 | 11.716801 | 9.252817 | -2.679634 | 2.774942 | 0.077631 |
| std | 11.971827 | 6.655333 | 9.861443 | 10.625164 | 8.800553 | 15.375478 |
| min | -25.824199 | -8.497562 | -23.666439 | -29.429655 | -22.033329 | -35.264019 |
| 25% | 0.231327 | 7.161564 | 2.648845 | -10.652694 | -4.098043 | -14.003670 |
| 50% | 12.754335 | 11.896021 | 11.422441 | -2.631318 | 2.484416 | 1.500838 |
| 75% | 17.364337 | 16.279210 | 16.503676 | 5.340314 | 9.660898 | 14.050512 |
| max | 32.268570 | 32.909917 | 31.230550 | 26.422798 | 29.312010 | 31.727042 |

```
In [15]: def visualize(ds_2, title = 'Clusters', n_components = 2):
    pca = PCA(n_components).fit_transform(ds_2)
    x = pca[:, 0]
    y = pca[:, 1]
    plt.scatter(x, y)
    plt.title(title)
    plt.xlabel('PCA 1')
    plt.ylabel('PCA 2')
```

```
plt.show()
visualize(ds_2)
```



```
In [16]: def kmean_silhouette(ds_2, to_print = False):
    best_score = -1
    scores = []
    clusters = [2, 3, 4, 5, 6, 7, 8, 9, 10, 20, 50, 100]

    ds_2_scaled = ds_2.copy()
    ds_2_scaled[ds_2_scaled.columns] = StandardScaler().fit_transform(ds_2_scaled)

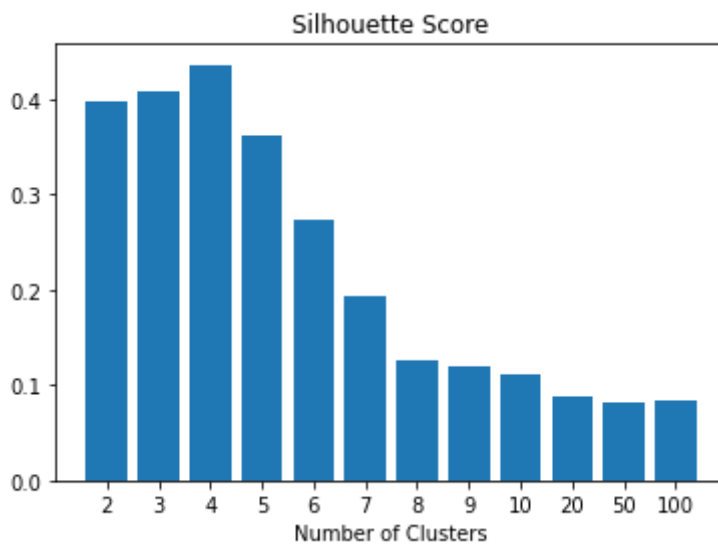
    for cluster in clusters:
        kmeans = KMeans(n_clusters = cluster).fit(ds_2_scaled)
        score = silhouette_score(ds_2_scaled, kmeans.labels_)
        scores.append(score)

        if (to_print): print('Cluster: {}, Score: {}'.format(cluster, score))

        if (score > best_score):
            best_score = score
            best_cluster = cluster

    plt.bar(range(len(scores)), list(scores), align = 'center')
    plt.xticks(range(len(scores)), list(clusters))
    plt.title('Silhouette Score')
    plt.xlabel('Number of Clusters')
    plt.show()
    return best_cluster

num_clusters = kmean_silhouette(ds_2, to_print = False)
print('Number of groups:', num_clusters)
```

Number of groups: 4

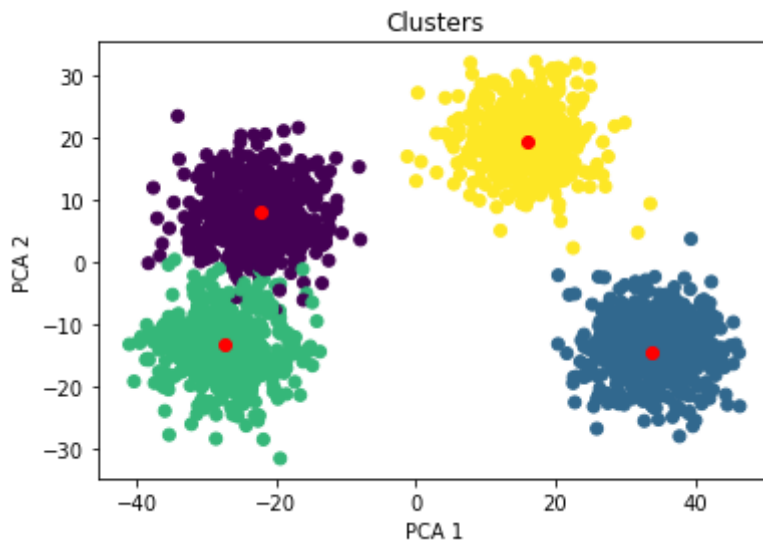
```
In [17]: def assign_kmeans(ds_2, num_clusters, n_components = 2):
    pca = PCA(n_components = n_components)
    kmeans = KMeans(n_clusters = num_clusters)
    kmeans.fit(ds_2)
    centers = kmeans.cluster_centers_
    centers_pca = pca.fit_transform(centers)
    return kmeans.labels_, centers_pca

labels, center_pca = assign_kmeans(ds_2, num_clusters)
print('Labels for each row:', labels)
```

Labels for each row: [0 3 3 ... 2 1 0]

```
In [18]: def visualizing_groups(ds_2, labels, center_pca, n_components = 2):
    pca = PCA(n_components = n_components).fit_transform(ds_2)
    x = pca[:, 0]
    y = pca[:, 1]
    plt.scatter(x, y, c = labels)
    plt.title('Clusters')
    plt.xlabel('PCA 1')
    plt.ylabel('PCA 2')
    plt.scatter(center_pca[:, 0], center_pca[:, 1], marker = 'o', color = 'red')
    plt.show()

visualizing_groups(ds_2, labels, center_pca)
```



Question 3

Stack Overflow provides a tool at <https://data.stackexchange.com/stackoverflow/query/new> that allows SQL queries to be run against their data. After reviewing the database schema provided on their site, please answer the questions below by providing both your answer and the query used to derive it.

a. How many posts were created in 2017

```
In [19]: sql = """SELECT COUNT(*) FROM Posts
            WHERE CreationDate BETWEEN '2017-01-01' AND '2017-12-31'
            """

# Note:
# usually you can use CreationDate::Date to get just the date value
# without the timestamp, but this server did not support it
```

Result: 5021226

b. What post/question received the most answers?

```
In [20]: sql = """SELECT TOP(1) Id, AnswerCount FROM Posts
            ORDER BY AnswerCount DESC
            """

# OR

sql = """SELECT Id, AnswerCount FROM Posts
            ORDER BY AnswerCount DESC
            LIMIT 1
            """

# Note:
# LIMIT not supported in this server, only TOP
```

Result:

Id: 184618518

AnswerCount: 518

c. For posts created in 2020, what were the top 10 tags?

```
In [21]: sql = """SELECT TOP(11) Tags, COUNT(*) FROM Posts
           WHERE CreationDate BETWEEN '2020-01-01' AND '2020-12-31'
           GROUP BY Tags
           ORDER BY COUNT(*) DESC
           """

# Note:
# I am showing TOP 11, becuse empty tag has the highest count, but not sure if
```

Result:

NO TAGS: 2442284

python: 16230

javascript: 12838

python, pandas: 9675

r: 8823

html, css: 7743

java: 7050

excel, vba: 6802

javascript, reactjs: 6717

c++: 6622

reactjs: 5463

d. *BONUS* For the questions created in 2017, what was the average time (in seconds) between when the question was created and when the accepted answer was provided?

```
In [22]: sql = """
           SELECT
             AVG(
               CAST(
                 DATEDIFF (SECOND, p.CreationDate, a.CreationDate)
                 AS BIGINT)
             )
           FROM
             Posts p
           INNER JOIN
             Posts a
           ON
             p.AcceptedAnswerId = a.Id
           WHERE p.CreationDate BETWEEN '2017-01-01' AND '2017-12-31'
           """
```

Result: 838231

In []: